

```
import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
```

```
!pip install pmdarima
```

```
Requirement already satisfied: pmdarima in /usr/local/lib/python3.7/dist-packages (1.8.
Requirement already satisfied: numpy>=1.19.3 in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.7/
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in /usr/local/lib/python3.7/d
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: Cython!=0.29.18,>=0.29 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.7/dist-pack
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: patsy>=0.5.2 in /usr/local/lib/python3.7/dist-packages (
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.7/dist-package
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dis
```


```
df = pd.read_csv('medical_time_series.csv')
print(df.head())
```

```
   Day  Revenue
0    1  0.000000
1    2 -0.292356
2    3 -0.327772
3    4 -0.339987
4    5 -0.124888
```

```
print(df.isnull().sum())
```

```
Day      0
Revenue  0
dtype: int64
```

```
plt.figure(figsize=(14,10))
plt.plot(df.Day, df.Revenue)
plt.ylabel('Price in Mils')
plt.xlabel('Day Number')
plt.show()
```



```
df.info
```

```
<bound method DataFrame.info of      Day    Revenue
0         1    0.000000
1         2   -0.292356
2         3   -0.327772
3         4   -0.339987
4         5   -0.124888
..      ...      ...
726    727   15.722056
727    728   15.865822
728    729   15.708988
729    730   15.822867
730    731   16.069429
```

```
[731 rows x 2 columns]>
```

```
print(df.duplicated().sum())
```

```
0
```

```
print(df.describe())
```


	Day	Revenue
count	731.000000	731.000000
mean	366.000000	14.179608
std	211.165812	6.959905
min	1.000000	-4.423299
25%	183.500000	11.121742
50%	366.000000	15.951830
75%	548.500000	19.293506
max	731.000000	24.792249

```
from statsmodels.tsa.stattools import adfuller
df = pd.read_csv('medical_time_series.csv')
new_result = adfuller(df.Revenue)
print('ADF Stat: %f' % new_result[0])
print('p-value: %f' % new_result[1])
print('Critical Values:')
for key, value in new_result[4].items():
    print('\t%s: %.3f' % (key, value))
```

```
ADF Stat: -2.218319
p-value: 0.199664
Critical Values:
    1%: -3.439
    5%: -2.866
   10%: -2.569
```

```
from sklearn.model_selection import train_test_split
train, test = train_test_split(df, test_size=0.2, random_state=42, shuffle=False)
```

```
plt.figure(figsize=(14,10))
plt.plot(train.Day, train.Revenue, color='blue')
plt.plot(test.Day, test.Revenue, color='red')
plt.ylabel('Price in Mils')
plt.xlabel('Day Number')
plt.show()
```



```
train.to_csv('train.csv')
test.to_csv('test.csv')
```

```
print(train.columns)
print(test.columns)
```

```
Index(['Day', 'Revenue'], dtype='object')
Index(['Day', 'Revenue'], dtype='object')
```

```
print(df)
```

	Day	Revenue
0	1	0.000000
1	2	-0.292356
2	3	-0.327772
3	4	-0.339987
4	5	-0.124888
..
726	727	15.722056
727	728	15.865822
728	729	15.708988
729	730	15.822867
730	731	16.069429

```
[731 rows x 2 columns]
```

```
print(df.columns)
#train_copy = train[['Day', 'Revenue']].copy()
#test_copy = test[['Day', 'Revenue']].copy()
df.to_numpy()
```

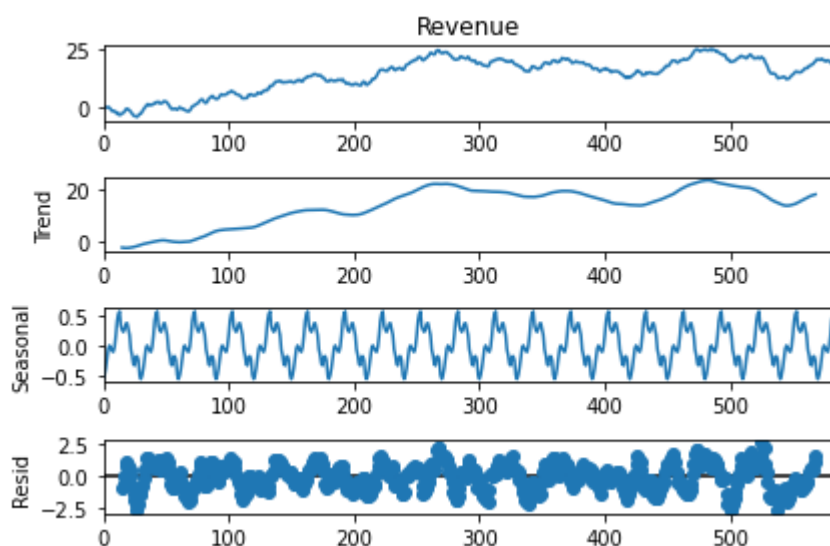
```
print(train.head)
print(test.head)
```

```
Index(['Day', 'Revenue'], dtype='object')
<bound method NDFrame.head of      Day    Revenue
0      1    0.000000
1      2   -0.292356
2      3   -0.327772
3      4   -0.339987
4      5   -0.124888
..    ..    ...
579   580   19.782635
580   581   19.088265
581   582   18.805501
582   583   18.910233
583   584   19.186089
```

```
[584 rows x 2 columns]>
<bound method NDFrame.head of      Day    Revenue
584   585   19.312734
585   586   19.576725
586   587   18.988035
587   588   18.437608
588   589   18.519085
..    ..    ...
726   727   15.722056
727   728   15.865822
728   729   15.708988
729   730   15.822867
730   731   16.069429
```


```
[147 rows x 2 columns]>
```

```
decomp = seasonal_decompose(train['Revenue'], model='additive', period=30) # freq or period.
decomp.plot()
plt.show()
```




```
decomp_trend = decomp.trend
decomp_trend.plot(figsize=(16,4))
plt.xlabel('Days')
plt.ylabel('Revenue in Mils')
```

```
Text(0, 0.5, 'Revenue in Mils')
```




```
from statsmodels.graphics import tsaplots
```

```
fig = tsaplots.plot_acf(df['Revenue'], lags = 30)
plt.show()
```



```
plt.psd(df['Revenue'])
```

```
(array([1.52612936e+04, 8.12118786e+03, 2.62760550e+02, 1.24454986e+02,
        5.40026379e+01, 2.55196420e+01, 2.89237553e+01, 6.29449800e+00,
        8.38967529e+00, 8.79735213e+00, 3.26244345e+00, 5.16464998e+00,
        4.64370619e+00, 1.96867736e+00, 1.48167320e+00, 2.11256649e+00,
        1.08550271e+00, 5.81988645e-01, 2.76744723e+00, 2.22308983e+00,
        5.50876230e-01, 7.49752606e-01, 1.26157210e+00, 8.77859381e-01,
        4.84404330e-01, 4.39565525e-01, 2.05055025e-01, 9.50457901e-01,
        5.23499207e-01, 4.35456489e-01, 1.69884060e+00, 8.97977002e-01,
        7.28275718e-01, 6.11527179e-01, 1.49053313e+00, 5.44383887e-01,
        2.17350902e-01, 5.07814496e-01, 4.43820257e-01, 2.94184611e-01,
        2.55373626e-01, 9.44702588e-02, 1.58619217e-01, 1.88829134e-01,
        2.11804293e-01, 6.04121179e-01, 4.95073852e-01, 2.15698285e-01,
        1.04432711e-01, 1.46393379e-01, 2.07457981e-01, 1.26602668e-01,
        1.74592665e-01, 1.42674366e-01, 9.63196690e-02, 1.17768273e-01,
        1.12081971e-02, 7.05631210e-02, 8.82393791e-02, 1.79463366e-02,
        4.63319518e-02, 8.22217635e-02, 8.62002792e-02, 8.11471442e-02,
        9.76265725e-02, 7.64686033e-02, 4.69018363e-02, 2.27885092e-02,
        3.22945840e-02, 4.15665768e-02, 3.87010238e-02, 1.37096984e-02,
        2.54563144e-02, 5.96461061e-02, 8.54027434e-02, 5.14038459e-02,
        3.17470361e-02, 2.66619122e-02, 5.20528881e-02, 3.74783288e-03,
        4.60111740e-02, 4.72280258e-02, 2.64401277e-02, 2.65654525e-02,
        2.74537510e-02, 1.30164715e-02, 2.52371618e-02, 5.09056679e-02,
        2.36007364e-02, 3.60978502e-02, 1.34740313e-02, 2.47166596e-02,
        5.63282623e-02, 7.25977232e-02, 2.73963882e-02, 3.40075699e-02,
        7.28076058e-03, 2.37376075e-02, 3.35477194e-02, 1.11675245e-02,
        3.86130031e-02, 5.52632330e-02, 3.06776062e-02, 1.02314465e-02,
        1.00437589e-02, 5.88763445e-02, 6.55634547e-02, 6.48669973e-02,
        4.56754349e-02, 1.17566185e-02, 4.41995523e-03, 2.14450712e-02,
        3.82815436e-02, 3.32170073e-02, 3.01496591e-02, 1.53250228e-02,
        5.43350164e-03, 1.84455310e-02, 1.19157543e-02, 3.05548420e-02,
        8.40529002e-02, 3.87920060e-02, 1.07052164e-02, 4.84092102e-02,
        4.49317201e-02, 3.20797981e-02, 2.10764097e-02, 4.37206681e-02,
        1.63676609e-02]),
array([0.      , 0.0078125, 0.015625 , 0.0234375, 0.03125  , 0.0390625,
        0.046875 , 0.0546875, 0.0625   , 0.0703125, 0.078125 , 0.0859375,
        0.09375  , 0.1015625, 0.109375 , 0.1171875, 0.125    , 0.1328125,
        0.140625 , 0.1484375, 0.15625  , 0.1640625, 0.171875 , 0.1796875,
        0.1875   , 0.1953125, 0.203125 , 0.2109375, 0.21875  , 0.2265625,
        0.234375 , 0.2421875, 0.25     , 0.2578125, 0.265625 , 0.2734375,
        0.28125  , 0.2890625, 0.296875 , 0.3046875, 0.3125   , 0.3203125,
        0.328125 , 0.3359375, 0.34375  , 0.3515625, 0.359375 , 0.3671875,
        0.375    , 0.3828125, 0.390625 , 0.3984375, 0.40625  , 0.4140625,
        0.421875 , 0.4296875, 0.4375   , 0.4453125, 0.453125 , 0.4609375,
        0.46875  , 0.4765625, 0.484375 , 0.4921875, 0.5      , 0.5078125,
        0.515625 , 0.5234375, 0.53125  , 0.5390625, 0.546875 , 0.5546875,
        0.5625   , 0.5703125, 0.578125 , 0.5859375, 0.59375  , 0.6015625,
        0.609375 , 0.6171875, 0.625    , 0.6328125, 0.640625 , 0.6484375,
        0.65625  , 0.6640625, 0.671875 , 0.6796875, 0.6875   , 0.6953125,
        0.703125 , 0.7109375, 0.71875  , 0.7265625, 0.734375 , 0.7421875,
        0.75     , 0.7578125, 0.765625 , 0.7734375, 0.78125  , 0.7890625,
        0.796875 , 0.8046875, 0.8125   , 0.8203125, 0.828125 , 0.8359375,
        0.84375  , 0.8515625, 0.859375 , 0.8671875, 0.875    , 0.8828125,
        0.890625 , 0.8984375, 0.90625  , 0.9140625, 0.921875 , 0.9296875,
        0.9375   , 0.9453125, 0.953125 , 0.9609375, 0.96875  , 0.9765625,
        0.984375 , 0.9921875, 1.      ]))
```



```
from pmdarima import auto_arima
auto_arima_model = auto_arima(df['Revenue'], start_p=1, start_q=1,
                               max_p=3, max_q=3, m=12,
                               start_P=0, seasonal=True,
                               d=1, D=1, trace=True,
                               error_action='ignore',
                               suppress_warnings=True,
                               stepwise=True)
print(auto_arima_model.aic())
print(auto_arima_model.summary())
```

Performing stepwise search to minimize aic

```
ARIMA(1,1,1)(0,1,1)[12]      : AIC=inf, Time=4.53 sec
ARIMA(0,1,0)(0,1,0)[12]      : AIC=1548.267, Time=0.08 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=1161.213, Time=0.54 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=inf, Time=7.12 sec
ARIMA(1,1,0)(0,1,0)[12]      : AIC=1413.401, Time=0.14 sec
ARIMA(1,1,0)(2,1,0)[12]      : AIC=1076.065, Time=1.72 sec
ARIMA(1,1,0)(2,1,1)[12]      : AIC=inf, Time=19.52 sec
ARIMA(1,1,0)(1,1,1)[12]      : AIC=inf, Time=9.33 sec
ARIMA(0,1,0)(2,1,0)[12]      : AIC=1229.919, Time=1.09 sec
ARIMA(2,1,0)(2,1,0)[12]      : AIC=1078.060, Time=2.29 sec
ARIMA(1,1,1)(2,1,0)[12]      : AIC=1078.060, Time=2.33 sec
ARIMA(0,1,1)(2,1,0)[12]      : AIC=1103.447, Time=0.97 sec
ARIMA(2,1,1)(2,1,0)[12]      : AIC=1078.936, Time=4.18 sec
ARIMA(1,1,0)(2,1,0)[12] intercept : AIC=1078.064, Time=3.69 sec
```

Best model: ARIMA(1,1,0)(2,1,0)[12]

Total fit time: 57.557 seconds

1076.0645264844293

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:
Model:          SARIMAX(1, 1, 0)x(2, 1, 0, 12)    Log Likelihood          -534.
Date:              Tue, 08 Mar 2022              AIC              1076.
Time:              00:02:36                      BIC              1094.
Sample:              0                          HQIC              1083.
```

```
              - 731
Covariance Type:      opg
```

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.4420      0.034      12.987      0.000      0.375      0.509
ar.S.L12       -0.7433      0.038     -19.718      0.000     -0.817     -0.669
ar.S.L24       -0.3460      0.038      -9.082      0.000     -0.421     -0.271
sigma2          0.2564      0.015      17.415      0.000      0.228      0.285
=====
```

```
Ljung-Box (L1) (Q):          0.01      Jarque-Bera (JB):          4.02
```


Prob(Q):	0.94	Prob(JB):	0.13
Heteroskedasticity (H):	0.96	Skew:	0.11
Prob(H) (two-sided):	0.75	Kurtosis:	2.70
=====			

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
#arima forecast
# SARIMAX(1, 1, 0)x(2, 1, 0, 12)
import statsmodels.api as sm
model = sm.tsa.statespace.SARIMAX(train['Revenue'], order=(1,1,0), seasonal_order = (2,1,0,12))
SARIMAXresults = model.fit()

#df['forecast'] = results.predict(start = 0, end = 730, dynamic = True)
#df['Revenue'].plot(figsize = 12,8)
result = SARIMAXresults.get_forecast()
forecast_test = test['Revenue'].values # using test set to see what predicted would be
forecast = result.predicted_mean
print('expected:', forecast)
print('forecast:', forecast_test[0])
print('std error:', result.se_mean)

expected: 584      19.740296
dtype: float64
forecast: 19.31273398
std error: 584      0.511004
dtype: float64

start = len(train.Revenue)
end = len(train.Revenue) + len(test.Revenue) - 1


plt.figure(figsize=(14,10))
predict = SARIMAXresults.predict(start, end, typ='levels')
plt.plot(train.Revenue, label = 'train set')
plt.plot(test.Revenue, label = 'test set')
plt.plot(predict, label = 'prediction')
plt.title('Revenue Predictions')
plt.xlabel('Day')
plt.xlabel('Revenue')
plt.legend(loc='upper left')
plt.show()
```



```
model = sm.tsa.statespace.SARIMAX(df['Revenue'], order=(1,1,0), seasonal_order=(2,1,0,12))
results = model.fit()
forecast = results.predict(start = len(df['Revenue']), end = (len(df['Revenue']) - 1) + 365)
```

```
df['Revenue'].plot()
forecast.plot(legend = True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f124d8cba50>



```
from sklearn.metrics import mean_squared_error
from statsmodels.tools.eval_measures import rmse
```

```
MSE = mean_squared_error(test['Revenue'], predict)
RMSE = rmse(test['Revenue'], predict)
print('MSE: ', MSE)
print('RMSE: ', RMSE)
```

```
MSE: 228.1822083849405
RMSE: 15.105701188125645
```

```
interval = [0.25, 0.1, 0.05, 0.01]
for i in interval:
    conf = result.conf_int(alpha=i)
    print("confidence intervals: ", ((1-i)*100), conf)
```

```
confidence intervals: 75.0      lower Revenue  upper Revenue
584      19.152463      20.32813
confidence intervals: 90.0      lower Revenue  upper Revenue
584      18.899769      20.580824
confidence intervals: 95.0      lower Revenue  upper Revenue
584      18.738746      20.741847
confidence intervals: 99.0      lower Revenue  upper Revenue
584      18.424036      21.056557
```

✓ 0s completed at 7:02 PM

